

Intro to {purrr}

And a more explicit description of functional programming

Daniel Anderson
Week 3, Class 1



Agenda

- Thinking more about functional programming
 - A small example
- Introduce you to {purrr} and contrast it with the base functions we learned last week

Learning objectives

- Understand how `purrr::map` relates to `lapply` and `for` loops
- Understand the six basic variants of `purrr::map`, when they should be used, and when they will fail
- Understand what functional programming is, and how `{purrr}` can help facilitate the process

Functional Programming

What is FP?

decomposing a big problem into smaller pieces, then solving each piece with a function or combination of functions

- Adv-R

Example

Calculate top mpg manufactures

- First, we'll subset the data to 4 cylinders. This will be the dataset we'll solve the problem on:

```
library(tidyverse)
four_cyl <- filter(mpg, cyl == 4)
```

- Next, we'll filter for cases where the city miles per gallon is in the top 10% (i.e., greater than or equal to the 90th percentile).

```
ninety <- four_cyl %>%  
  filter(cty >= quantile(cty, probs = 0.9))  
ninety
```

```
## # A tibble: 10 x 11  
##   manufacturer model displ  year   cyl trans  drv    cty   hwy fl   class  
##   <chr>         <chr> <dbl> <int> <int> <chr> <chr> <int> <int> <chr> <chr>  
## 1 honda        civic  1.6  1999     4 manu... f      28    33 r    subc...  
## 2 honda        civic  1.6  1999     4 manu... f      25    32 r    subc...  
## 3 honda        civic  1.8  2008     4 manu... f      26    34 r    subc...  
## 4 honda        civic  1.8  2008     4 auto... f      25    36 r    subc...  
## 5 toyota       coro... 1.8  1999     4 manu... f      26    35 r    comp...  
## 6 toyota       coro... 1.8  2008     4 manu... f      28    37 r    comp...  
## 7 toyota       coro... 1.8  2008     4 auto... f      26    35 r    comp...  
## 8 volkswagen   jetta  1.9  1999     4 manu... f      33    44 d    comp...  
## 9 volkswagen   new ... 1.9  1999     4 manu... f      35    44 d    subc...  
## 10 volkswagen  new ... 1.9  1999     4 auto... f      29    41 d    subc...
```

-
- Now count the unique occurrences for each manufacturer, manufacturer/model, and class

```
count_manufacturer <- count(ninety, manufacturer)
count_model <- count(ninety, manufacturer, model)
count_class <- count(ninety, class)
```


- Produce a plot for each

```
plot_manufacturer <-  
  ggplot(count_manufacturer, aes(fct_reorder(manufacturer, -n), n)) +  
    geom_col(aes(fill = manufacturer)) +  
    scale_fill_brewer(palette = "Set3") +  
    labs(title = "Manufacturers",  
         x = "",  
         y = "") +  
    guides(fill = "none")
```

```
plot_car <- count_model %>%  
  unite(car, manufacturer, model, sep = " ") %>%  
  ggplot(aes(fct_reorder(car, -n), n)) +  
    geom_col(aes(fill = car)) +  
    scale_fill_brewer(palette = "Set3") +  
    labs(title = "Top 10% of city mpg",  
         subtitle = "Car frequency",  
         x = "",  
         y = "") +  
    guides(fill = "none")
```


```
plot_class <-  
  ggplot(count_class, aes(fct_reorder(class, -n), n)) +  
    geom_col(aes(fill = class)) +  
    scale_fill_brewer(palette = "Set3") +  
    labs(title = "Car Class",
```

Assemble the plots


```
library(patchwork)  
plot_car / (plot_manufacturer + plot_class)
```

Top 10% of city mpg


Car frequency



Manufacturers



Car Class



Functional Programming Version

At least in spirit

Filter all

```
by_cyl <- split(mpg, mpg$cyl)
top_10 <- lapply(by_cyl, function(x) {
  filter(x, cty >= quantile(cty, probs = 0.9))
})
str(top_10)
```

```
## List of 4
## $ 4:Classes 'tbl_df', 'tbl' and 'data.frame': 10 obs. of 11 variables:
## ..$ manufacturer: chr [1:10] "honda" "honda" "honda" "honda" ...
## ..$ model       : chr [1:10] "civic" "civic" "civic" "civic" ...
## ..$ displ       : num [1:10] 1.6 1.6 1.8 1.8 1.8 1.8 1.8 1.9 1.9 1.9
## ..$ year        : int [1:10] 1999 1999 2008 2008 1999 2008 2008 1999 1999 199
## ..$ cyl         : int [1:10] 4 4 4 4 4 4 4 4 4 4
## ..$ trans       : chr [1:10] "manual(m5)" "manual(m5)" "manual(m5)" "auto(l5)
## ..$ drv         : chr [1:10] "f" "f" "f" "f" ...
## ..$ cty         : int [1:10] 28 25 26 25 26 28 26 33 35 29
## ..$ hwy         : int [1:10] 33 32 34 36 35 37 35 44 44 41
## ..$ fl         : chr [1:10] "r" "r" "r" "r" ...
## ..$ class       : chr [1:10] "subcompact" "subcompact" "subcompact" "subcompa
## $ 5:Classes 'tbl_df', 'tbl' and 'data.frame': 2 obs. of 11 variables:
## ..$ manufacturer: chr [1:2] "volkswagen" "volkswagen"
## ..$ model       : chr [1:2] "jetta" "jetta"
## ..$ displ       : num [1:2] 2.5 2.5
## ..$ year        : int [1:2] 2008 2008
```

All counts

```
counts <- lapply(top_10, function(x) {  
  count_manufacturer <- count(x, manufacturer)  
  count_class <- count(x, class)  
  
  count_model <- count(x, manufacturer, model) %>%  
    unite(car, manufacturer, model, sep = " ")  
  
  return(list(mfr = count_manufacturer,  
             car = count_model,  
             class = count_class))  
})  
str(counts)
```

```
## List of 4  
## $ 4:List of 3  
## ..$ mfr :Classes 'tbl_df', 'tbl' and 'data.frame': 3 obs. of 2 variables  
## .. ..$ manufacturer: chr [1:3] "honda" "toyota" "volkswagen"  
## .. ..$ n : int [1:3] 4 3 3  
## ..$ car :Classes 'tbl_df', 'tbl' and 'data.frame': 4 obs. of 2 variables  
## .. ..$ car: chr [1:4] "honda civic" "toyota corolla" "volkswagen jetta" "volk  
## .. ..$ n : int [1:4] 4 3 1 2  
## ..$ class:Classes 'tbl_df', 'tbl' and 'data.frame': 2 obs. of 2 variables  
## .. ..$ class: chr [1:2] "compact" "subcompact"
```

Plots

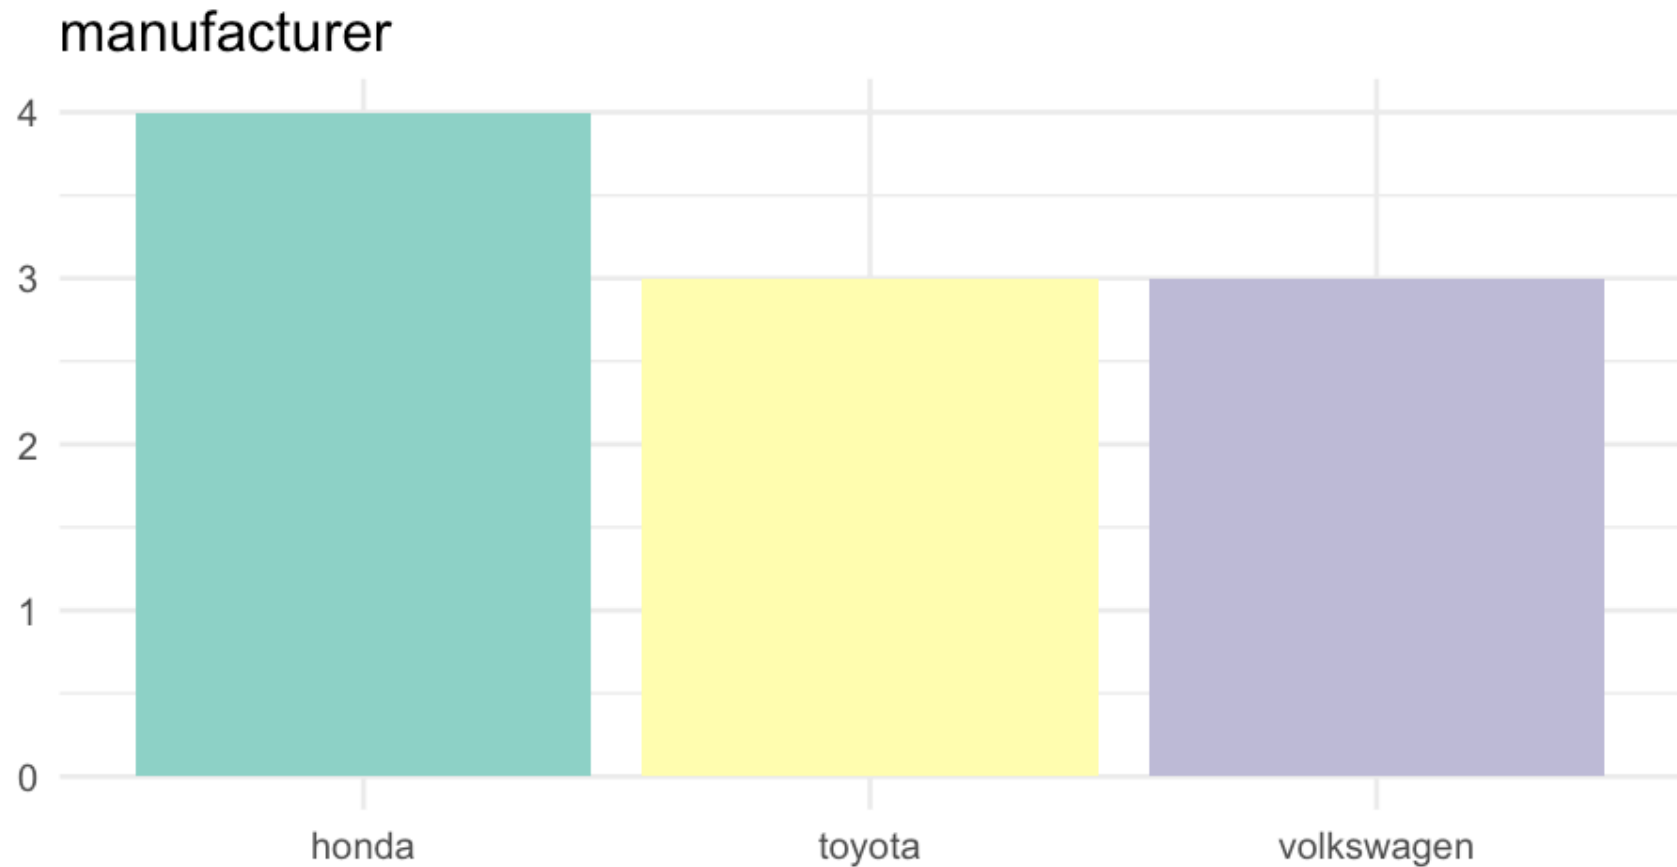
Let's write a couple functions

(I recognize we haven't talked about functions yet, but as a preview...)


```
indiv_plot <- function(d, var) {  
  p <- ggplot(d, aes_string(fct_reorder(d[[var]], -d[["n"]]),  
                             d[["n"]])) +  
    geom_col(aes_string(fill = var)) +  
    scale_fill_brewer(palette = "Set3") +  
    labs(title = var,  
         x = "",  
         y = "") +  
    guides(fill = "none")  
  
  if(var == "car") {  
    p <- p + labs(title = "Top 10% of city mpg",  
                  subtitle = var)  
  }  
  return(p)  
}
```

Test it

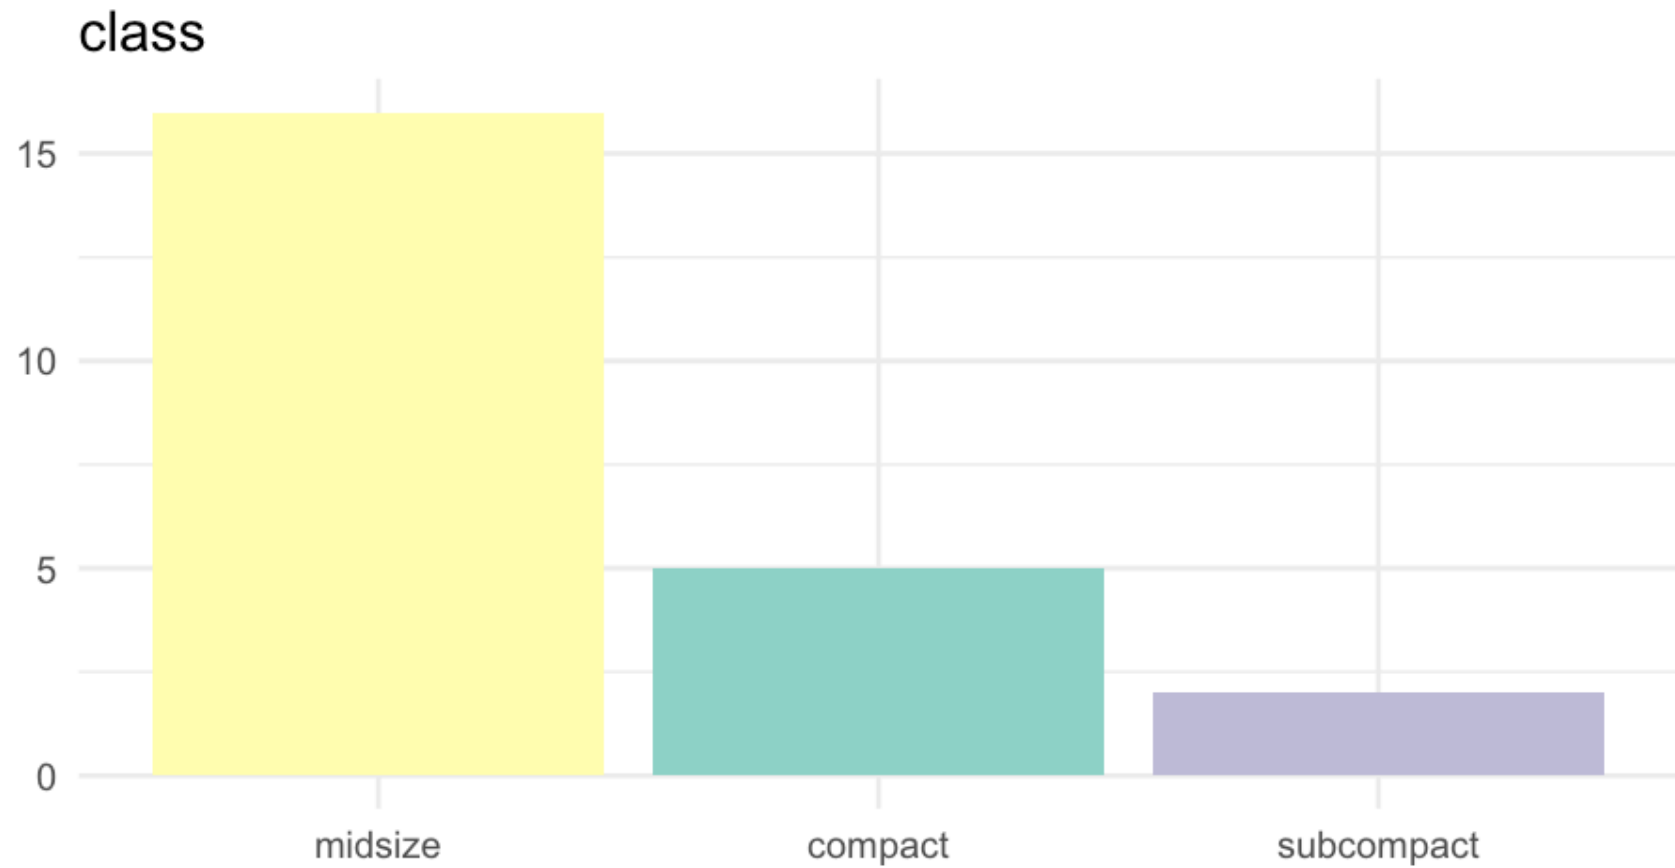
```
indiv_plot(counts[["4"]]  
$mfr, "manufacturer")
```



```
indiv_plot(counts[["4"]] $\$$ car, "car")
```




```
indiv_plot(counts[["6"]] $\$$ class, "class")
```



Compile plots function


```
full_plot <- function(l) {  
  indiv_plot(l[["car"]], "car") / (  
    indiv_plot(l[["mfr"]], "manufacturer") +  
    indiv_plot(l[["class"]], "class")  
  )  
}
```

Test it


```
full_plot(counts[[1]])
```

Top 10% of city mpg


car



manufacturer



class




Finish up


```
plots <- lapply(counts, full_plot)
```

```
plots[[1]]
```


Top 10% of city mpg
car



manufacturer



class



```
plots[[2]]
```


Top 10% of city mpg
car



manufacturer




class




```
plots[[3]]
```


Top 10% of city mpg
car



manufacturer




class




```
plots[[4]]
```


Top 10% of city mpg car



manufacturer



class



{purrr}

functionals

| a function that takes a function as input and returns a vector as output.

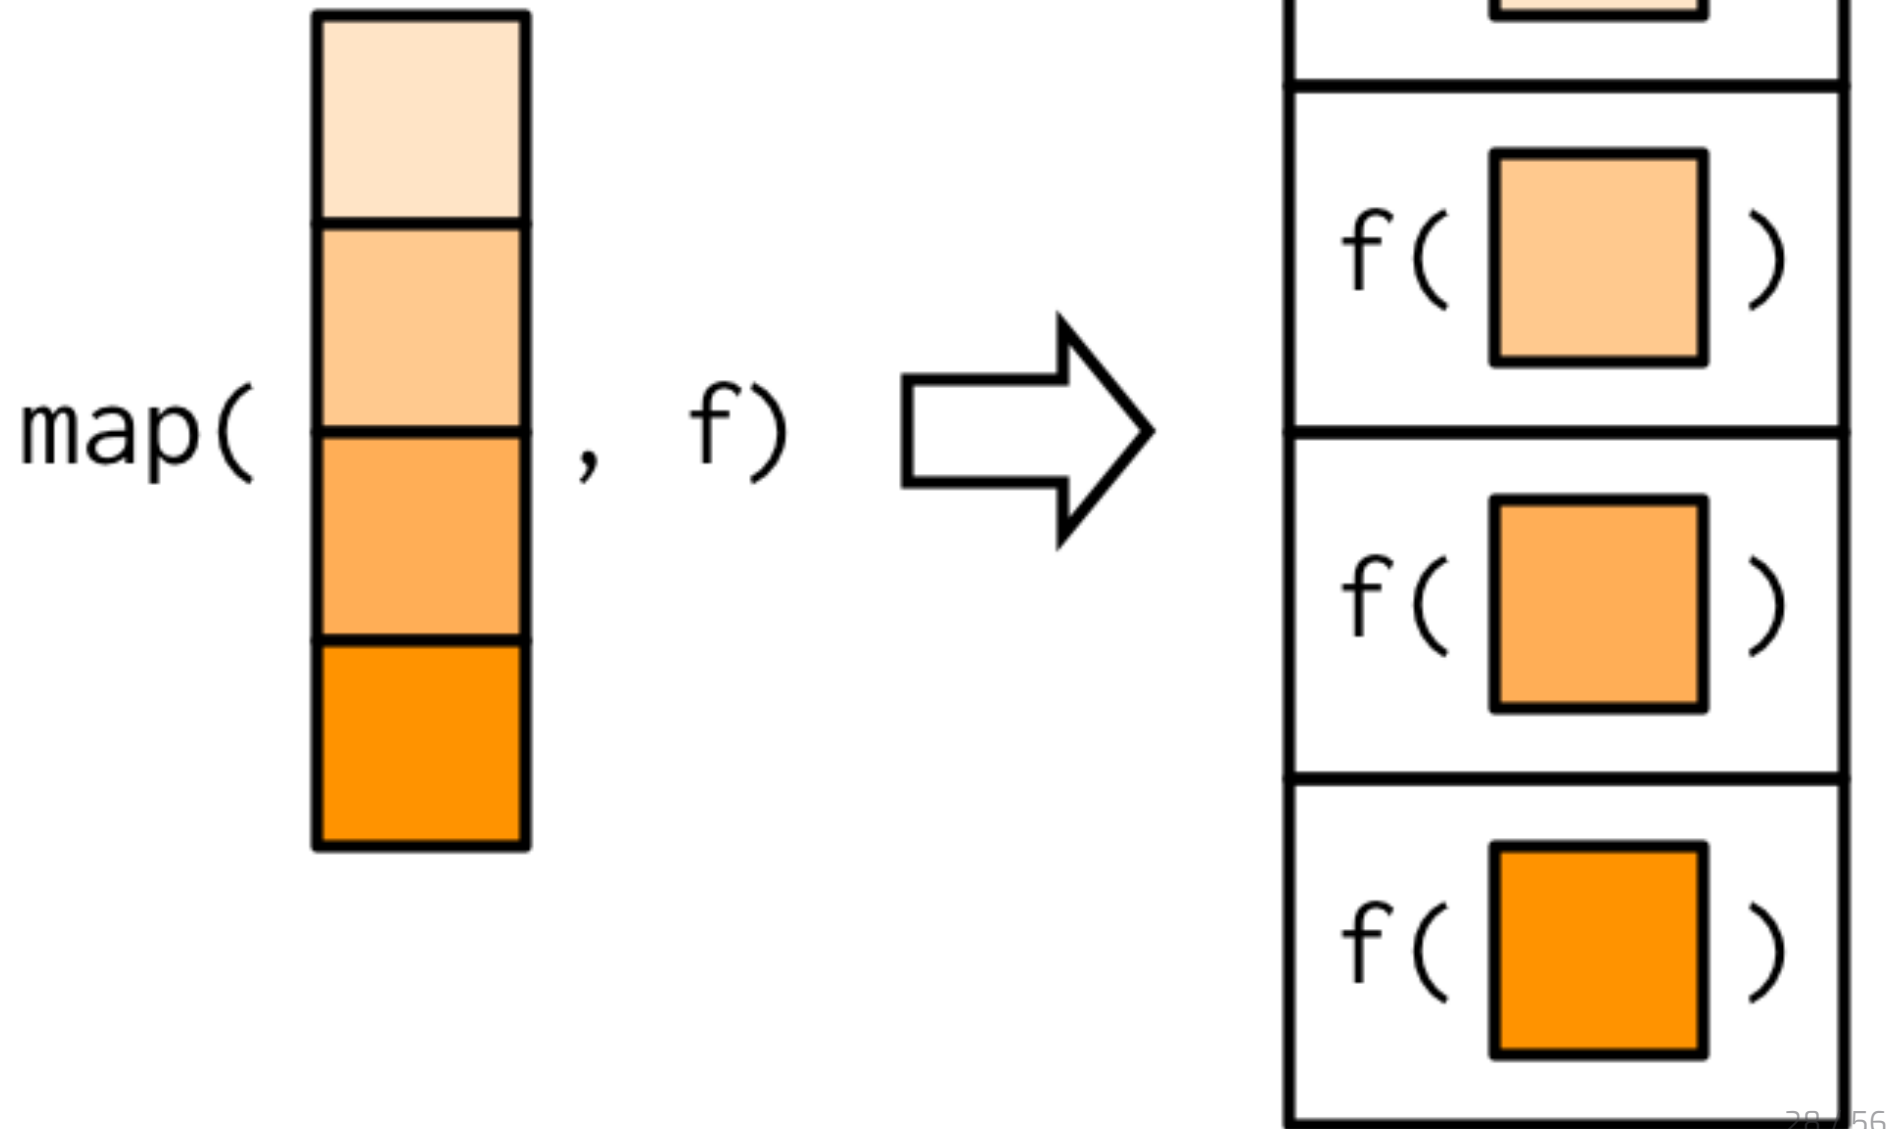
- What does this mean?

purrr::map

```
library(purrr) # loaded automatically by tidyverse  
map(1:3, rnorm)
```

```
## [[1]]  
## [1] 0.5941747  
##  
## [[2]]  
## [1] 0.5532272 0.2522552  
##  
## [[3]]  
## [1] 0.6873934 -0.1936925 -0.2495917
```

Graphically



Comparison to base::lapply

lapply

```
lapply(1:3, rnorm)
```

```
## [[1]]  
## [1] -0.7838051  
##  
## [[2]]  
## [1] -1.7543676  0.0689319  
##  
## [[3]]  
## [1]  0.07790071 -0.65711040 -0.12277615
```

map

```
map(1:3, rnorm)
```

```
## [[1]]  
## [1] -0.9771206  
##  
## [[2]]  
## [1] -0.4845289  1.9154039  
##  
## [[3]]  
## [1] -1.36694488  0.04390757 -0.07050604
```

side note: What exactly is going on here?

The base equivalent to `map()` is `lapply()`. The only difference is that `lapply()` does not support the helpers that you'll learn about [next], so if you're only using `map()` from `purrr`, you can skip the additional dependency and use `lapply()` directly.

- Adv R

Equivalents

The following are equivalent

```
map(mtcars, function(x) length(unique(x)))
```

```
lapply(mtcars, function(x) length(unique(x)))
```

Equivalents

The following are equivalent

```
map(mtcars, function(x) length(unique(x)))  
lapply(mtcars, function(x) length(unique(x)))
```

{purrr} also allows you to specify anonymous functions more succinctly using the formula interface

```
map(mtcars, ~length(unique(.x)))
```

```
## $mpg  
## [1] 25  
##  
## $cyl  
## [1] 3  
##  
## $disp  
## [1] 27  
##  
## $hp  
## [1] 22
```

Extracting elements

```
l <- list(  
  list(-1, x = 1, y = c(2), z = "a"),  
  list(-2, x = 4, y = c(5, 6), z = "b"),  
  list(-3, x = 8, y = c(9, 10, 11))  
)
```

Extracting elements

```
l <- list(  
  list(-1, x = 1, y = c(2), z = "a"),  
  list(-2, x = 4, y = c(5, 6), z = "b"),  
  list(-3, x = 8, y = c(9, 10, 11))  
)
```

Extract second element from each

```
map(l, 2)
```

```
## [[1]]  
## [1] 1  
##  
## [[2]]  
## [1] 4  
##  
## [[3]]  
## [1] 8
```

Doesn't work for `lapply`

```
lapply(l, 2)
```

```
## Error in match.fun(FUN): '2' is not a function, character or symbol
```

Doesn't work for `lapply`

```
lapply(l, 2)
```

```
## Error in match.fun(FUN): '2' is not a function, character or symbol
```

Instead, you have to apply an anonymous function

```
lapply(l, function(x) x[[2]])
```

```
## [[1]]  
## [1] 1  
##  
## [[2]]  
## [1] 4  
##  
## [[3]]  
## [1] 8
```

Alternatively the following is also the same

```
lapply(l, `[`, 2)
```

```
## [[1]]  
## [1] 1  
##  
## [[2]]  
## [1] 4  
##  
## [[3]]  
## [1] 8
```

Extract by name

```
map(l, "y")
```

```
## [[1]]  
## [1] 2  
##  
## [[2]]  
## [1] 5 6  
##  
## [[3]]  
## [1] 9 10 11
```

Multiple arguments

```
map(l, list("y", 1))
```

```
## [[1]]  
## [1] 2  
##  
## [[2]]  
## [1] 5  
##  
## [[3]]  
## [1] 9
```

{purrr} variants

Return a vector

- `map_dbl`
- `map_int`
- `map_char`
- `map_lgl`

```
str(l)
```

```
## List of 3
## $ :List of 4
## ..$ : num -1
## ..$ x: num 1
## ..$ y: num 2
## ..$ z: chr "a"
## $ :List of 4
## ..$ : num -2
## ..$ x: num 4
## ..$ y: num [1:2] 5 6
## ..$ z: chr "b"
## $ :List of 3
## ..$ : num -3
## ..$ x: num 8
## ..$ y: num [1:3] 9 10 11
```

```
str(l)
```

```
## List of 3
## $ :List of 4
## ..$ : num -1
## ..$ x: num 1
## ..$ y: num 2
## ..$ z: chr "a"
## $ :List of 4
## ..$ : num -2
## ..$ x: num 4
## ..$ y: num [1:2] 5 6
## ..$ z: chr "b"
## $ :List of 3
## ..$ : num -3
## ..$ x: num 8
## ..$ y: num [1:3] 9 10 11
```

```
map_dbl(l, "x")
```

```
## [1] 1 4 8
```

```
map_dbl(l, 1)
```

```
## [1] -1 -2 -3
```

Type match

- Coercion will occur if you request a different type

```
map_chr(l, "x")
```

```
## [1] "1.000000" "4.000000" "8.000000"
```

Type match

- Coercion will occur if you request a different type

```
map_chr(l, "x")
```

```
## [1] "1.000000" "4.000000" "8.000000"
```

- You'll get an error if element doesn't exist

```
map_chr(l, "z")
```

```
## Error: Result 3 must be a single string, not NULL of length 0
```

```
## Backtrace:
```

```
## [90m      [39m■
```

```
## [90m  1. [39m└─rmarkdown::render(...)
```

```
## [90m  2. [39m└─knitr::knit(...)
```

```
## [90m  3. [39m└─knitr::process_file(text, output)
```

```
## [90m  4. [39m└─base::withCallingHandlers(...)
```

```
## [90m  5. [39m└─knitr::process_group(group)
```

```
## [90m  6. [39m└─knitr::process_group.block(group)
```

```
## [90m  7. [39m└─knitr::call_block(x)
```

```
## [90m  8. [39m└─knitr::block_exec(params)
```

-
- Unless you set a default value

```
map_chr(l, "z", .default = NA_character_)
```

```
## [1] "a" "b" NA
```

Quick note on missing values

- In the prior case, specifying `NA` would work, instead of `NA_character_`

Quick note on missing values

- In the prior case, specifying `NA` would work, instead of `NA_character_`
- Generally, I think it's better to specify the type.

Quick note on missing values

- In the prior case, specifying `NA` would work, instead of `NA_character_`
- Generally, I think it's better to specify the type.
- **General programming rule:** The more strict the better (leads to fewer unexpected results)

Quick note on missing values

- In the prior case, specifying `NA` would work, instead of `NA_character_`
- Generally, I think it's better to specify the type.
- **General programming rule:** The more strict the better (leads to fewer unexpected results)
- Because (base) R likes to be inconsistent, here are they `NA` types

Type	NA value
character	<code>NA_character_</code>
integer	<code>NA_integer_</code>
double	<code>NA_real_</code>
logical	<code>NA</code> (see here)

Quick note on missing values

- In the prior case, specifying `NA` would work, instead of `NA_character_`
- Generally, I think it's better to specify the type.
- **General programming rule:** The more strict the better (leads to fewer unexpected results)
- Because (base) R likes to be inconsistent, here are the `NA` types

Type	NA value
character	<code>NA_character_</code>
integer	<code>NA_integer_</code>
double	<code>NA_real_</code>
logical	<code>NA</code> (see here)

```
typeof(NA)
```

```
## [1] "logical"
```

Some examples

Please follow along

```
econ <- economics %>%  
  mutate(year = lubridate::year(date))  
econ
```

```
## # A tibble: 574 x 7  
##   date           pce    pop psavert  uempmed  unemploy  year  
##   <date>         <dbl> <int>   <dbl>    <dbl>    <int> <dbl>  
## 1 1967-07-01  507.4   198712    12.5  4.5      2944  1967  
## 2 1967-08-01  510.5   198911    12.5  4.7      2945  1967  
## 3 1967-09-01  516.3000 199113    11.7  4.600000 2958  1967  
## 4 1967-10-01  512.9   199311    12.5  4.9      3143  1967  
## 5 1967-11-01  518.1   199498    12.5  4.7      3066  1967  
## 6 1967-12-01  525.800 199657    12.1  4.8      3018  1967  
## 7 1968-01-01  531.5   199808    11.7  5.100000 2878  1968  
## 8 1968-02-01  534.2   199920    12.2  4.5      3001  1968  
## 9 1968-03-01  544.9   200056    11.6  4.100000 2877  1968  
## 10 1968-04-01 544.6   200208    12.2  4.600000 2709  1968  
## # ... with 564 more rows
```

by_year

```
by_year <- split(econ, econ$year)
str(by_year)
```

```
## List of 49
## $ 1967:Classes 'tbl_df', 'tbl' and 'data.frame': 6 obs. of 7 variables:
## ..$ date : Date[1:6], format: "1967-07-01" ...
## ..$ pce : num [1:6] 507 510 516 513 518 ...
## ..$ pop : int [1:6] 198712 198911 199113 199311 199498 199657
## ..$ psavert : num [1:6] 12.5 12.5 11.7 12.5 12.5 12.1
## ..$ uempmed : num [1:6] 4.5 4.7 4.6 4.9 4.7 4.8
## ..$ unemploy: int [1:6] 2944 2945 2958 3143 3066 3018
## ..$ year : num [1:6] 1967 1967 1967 1967 1967 ...
## $ 1968:Classes 'tbl_df', 'tbl' and 'data.frame': 12 obs. of 7 variables:
## ..$ date : Date[1:12], format: "1968-01-01" ...
## ..$ pce : num [1:12] 532 534 545 545 550 ...
## ..$ pop : int [1:12] 199808 199920 200056 200208 200361 200536 200706 200
## ..$ psavert : num [1:12] 11.7 12.2 11.6 12.2 12 11.6 10.6 10.4 10.4 10.6 ...
## ..$ uempmed : num [1:12] 5.1 4.5 4.1 4.6 4.4 4.4 4.5 4.2 4.6 4.8 ...
## ..$ unemploy: int [1:12] 2878 3001 2877 2709 2740 2938 2883 2768 2686 2689 ..
## ..$ year : num [1:12] 1968 1968 1968 1968 1968 ...
## $ 1969:Classes 'tbl_df', 'tbl' and 'data.frame': 12 obs. of 7 variables:
## ..$ date : Date[1:12], format: "1969-01-01" ...
## ..$ pce : num [1:12] 584 590 590 595 601 ...
## ..$ pop : int [1:12] 201760 201881 202023 202161 202331 202507 202677 202
```

Fit a simple model to each year

Notes:

- We'll discuss a more elegant way to do this later
- This is not (statistically) the best way to approach this problem
- It's a good illustration, and in my experience there are lots of times with this approach works well, even if this particular example is a bit artificial

What is the relation between personal consumption expenditures (**pce**) and the unemployment percentage?

What is the relation between personal consumption expenditures (**pce**) and the unemployment percentage?

Problem: We don't have the percentage. Let's compute!

You try first!


```
perc <- map(by_year, ~mutate(.x, percent = unemploy / pop))
str(perc)
```

```
## List of 49
## $ 1967:Classes 'tbl_df', 'tbl' and 'data.frame': 6 obs. of 8 variables:
## ..$ date      : Date[1:6], format: "1967-07-01" ...
## ..$ pce       : num [1:6] 507 510 516 513 518 ...
## ..$ pop       : int [1:6] 198712 198911 199113 199311 199498 199657
## ..$ psavert   : num [1:6] 12.5 12.5 11.7 12.5 12.5 12.1
## ..$ uempmed   : num [1:6] 4.5 4.7 4.6 4.9 4.7 4.8
## ..$ unemploy  : int [1:6] 2944 2945 2958 3143 3066 3018
## ..$ year      : num [1:6] 1967 1967 1967 1967 1967 ...
## ..$ percent   : num [1:6] 0.0148 0.0148 0.0149 0.0158 0.0154 ...
## $ 1968:Classes 'tbl_df', 'tbl' and 'data.frame': 12 obs. of 8 variables:
## ..$ date      : Date[1:12], format: "1968-01-01" ...
## ..$ pce       : num [1:12] 532 534 545 545 550 ...
## ..$ pop       : int [1:12] 199808 199920 200056 200208 200361 200536 200706 200
## ..$ psavert   : num [1:12] 11.7 12.2 11.6 12.2 12 11.6 10.6 10.4 10.4 10.6 ...
## ..$ uempmed   : num [1:12] 5.1 4.5 4.1 4.6 4.4 4.4 4.5 4.2 4.6 4.8 ...
## ..$ unemploy  : int [1:12] 2878 3001 2877 2709 2740 2938 2883 2768 2686 2689 ..
## ..$ year      : num [1:12] 1968 1968 1968 1968 1968 ...
## ..$ percent   : num [1:12] 0.0144 0.015 0.0144 0.0135 0.0137 ...
## $ 1969:Classes 'tbl_df', 'tbl' and 'data.frame': 12 obs. of 8 variables:
## ..$ date      : Date[1:12], format: "1969-01-01" ...
## ..$ pce       : num [1:12] 584 590 590 595 601 ...
## ..$ pop       : int [1:12] 201760 201881 202023 202161 202331 202507 202677 202
## ..$ psavert   : num [1:12] 10 9.4 9.9 9.5 10 10.9 11.7 11.5 11.5 11.3 ...47 / 56
```

Fit the models

Fit a model of the form `lm(percent ~ pce)` to each year

You try first!

Fit the models

Fit a model of the form `lm(percent ~ pce)` to each year

You try first!

```
mods <- map(perc, ~lm(percent ~ pce, data = .x))
str(mods)
```

```
## List of 49
## $ 1967:List of 12
## ..$ coefficients : Named num [1:2] 8.39e-03 1.31e-05
## .. ..- attr(*, "names")= chr [1:2] "(Intercept)" "pce"
## ..$ residuals : Named num [1:6] -0.000205 -0.000255 -0.000281 0.000677 0.0
## .. ..- attr(*, "names")= chr [1:6] "1" "2" "3" "4" ...
## ..$ effects : Named num [1:6] -0.037041 0.00019 -0.000261 0.000737 0.000
## .. ..- attr(*, "names")= chr [1:6] "(Intercept)" "pce" "" "" ...
## ..$ rank : int 2
## ..$ fitted.values: Named num [1:6] 0.015 0.0151 0.0151 0.0151 0.0152 ...
## .. ..- attr(*, "names")= chr [1:6] "1" "2" "3" "4" ...
## ..$ assign : int [1:2] 0 1
## ..$ qr :List of 5
## .. ..$ qr : num [1:6, 1:2] -2.449 0.408 0.408 0.408 0.408 ...
## .. .. ..- attr(*, "dimnames")=List of 2
```

Extract coefficients

You try first

Hint: use `coef`. For example, see `coef(mods[[1]])`

Extract coefficients

You try first

Hint: use `coef`. For example, see `coef(mods[[1]])`

```
coefs <- map(mods, coef)
coefs[c(1:2, length(coefs))]
```

```
## $`1967`
##   (Intercept)          pce
## 8.388042e-03 1.307101e-05
##
## $`1968`
##   (Intercept)          pce
## 2.785458e-02 -2.495714e-05
##
## $`2015`
##   (Intercept)          pce
## 1.908717e-01 -1.350436e-05
```

Extract slopes

AKA - the coefficient that is not the intercept

You try first

Extract slopes

AKA - the coefficient that is not the intercept

You try first


```
slopes <- map_dbl(coefs, 2)
slopes
```

```
##           1967           1968           1969           1970           1971
## 1.307101e-05 -2.495714e-05 3.188337e-05 2.130891e-04 1.176607e-05
##           1972           1973           1974           1975           1976
## -2.694906e-05 -5.988616e-06 8.064082e-05 -1.529659e-06 9.390419e-06
##           1977           1978           1979           1980           1981
## -3.343402e-05 -1.149169e-05 4.999290e-06 2.361504e-05 3.211201e-05
##           1982           1983           1984           1985           1986
## 6.278562e-05 -4.447567e-05 -1.688549e-05 -6.769360e-06 -6.697402e-06
##           1987           1988           1989           1990           1991
## -1.959304e-05 -7.205655e-06 4.437364e-06 2.609293e-05 1.842054e-05
##           1992           1993           1994           1995           1996
## 6.051336e-07 -1.369550e-05 -2.092534e-05 7.668013e-07 -5.571115e-06
##           1997           1998           1999           2000           2001
## -9.922250e-06 -2.262969e-06 -2.872153e-06 -1.595609e-06 2.935469e-05
##           2002           2003           2004           2005           2006
```

Plot

- I trust you can do this

```
relation <- tibble(year = names(slopes),  
                   slope = slopes)  
  
ggplot(relation, aes(slope)) +  
  geom_histogram(fill = "cornflowerblue",  
                 color = "white")
```



Piping

We could also have done the previous in a pipeline.

```
by_year %>%  
  map(~mutate(.x, percent = unemploy / pop))
```

```
## $`1967`  
## # A tibble: 6 x 8  
##   date          pce    pop psavert  uempmed unemploy  year    percent  
##   <date>        <dbl> <int>   <dbl>   <dbl>   <int> <dbl>    <dbl>  
## 1 1967-07-01  507.4   198712   12.5  4.5     2944  1967  0.01481541  
## 2 1967-08-01  510.5   198911   12.5  4.7     2945  1967  0.01480562  
## 3 1967-09-01  516.3000 199113   11.7  4.600000 2958  1967  0.01485589  
## 4 1967-10-01  512.9   199311   12.5  4.9     3143  1967  0.01576933  
## 5 1967-11-01  518.1   199498   12.5  4.7     3066  1967  0.01536858  
## 6 1967-12-01  525.800 199657   12.1  4.8     3018  1967  0.01511592  
##  
## $`1968`  
## # A tibble: 12 x 8  
##   date          pce    pop psavert  uempmed unemploy  year    percent  
##   <date>        <dbl> <int>   <dbl>   <dbl>   <int> <dbl>    <dbl>  
## 1 1968-01-01  531.5   199808   11.7  5.100000 2878  1968  0.01440383  
## 2 1968-02-01  534.2   199920   12.2  4.5     3001  1968  0.01501100  
## 3 1968-03-01  544.9   200056   11.6  4.100000 2877  1968  0.01438097
```

```
by_year %>%  
  map(~mutate(.x, percent = unemploy / pop)) %>%  
  map(~lm(percent ~ pce, data = .x))
```

```
## `$1967`  
##  
## Call:  
## lm(formula = percent ~ pce, data = .x)  
##  
## Coefficients:  
## (Intercept)          pce  
## 8.388e-03      1.307e-05  
##  
##  
## `$1968`  
##  
## Call:  
## lm(formula = percent ~ pce, data = .x)  
##  
## Coefficients:  
## (Intercept)          pce  
## 2.785e-02     -2.496e-05  
##  
##  
## `$1969`  
##  
## Call:
```

```
by_year %>%  
  map(~mutate(.x, percent = unemploy / pop)) %>%  
  map(~lm(percent ~ pce, data = .x)) %>%  
  map(coef)
```

```
## $`1967`  
## (Intercept)          pce  
## 8.388042e-03 1.307101e-05  
##  
## $`1968`  
## (Intercept)          pce  
## 2.785458e-02 -2.495714e-05  
##  
## $`1969`  
## (Intercept)          pce  
## -5.308326e-03 3.188337e-05  
##  
## $`1970`  
## (Intercept)          pce  
## -0.1178905392 0.0002130891  
##  
## $`1971`  
## (Intercept)          pce  
## 1.594169e-02 1.176607e-05  
##  
## $`1972`  
## (Intercept)          pce
```

```
slopes <- by_year %>%
  map(~mutate(.x, percent = unemploy / pop)) %>%
  map(~lm(percent ~ pce, data = .x)) %>%
  map(coef) %>%
  map_dbl(2)
slopes
```

```
##           1967           1968           1969           1970           1971
## 1.307101e-05 -2.495714e-05 3.188337e-05 2.130891e-04 1.176607e-05
##           1972           1973           1974           1975           1976
## -2.694906e-05 -5.988616e-06 8.064082e-05 -1.529659e-06 9.390419e-06
##           1977           1978           1979           1980           1981
## -3.343402e-05 -1.149169e-05 4.999290e-06 2.361504e-05 3.211201e-05
##           1982           1983           1984           1985           1986
## 6.278562e-05 -4.447567e-05 -1.688549e-05 -6.769360e-06 -6.697402e-06
##           1987           1988           1989           1990           1991
## -1.959304e-05 -7.205655e-06 4.437364e-06 2.609293e-05 1.842054e-05
##           1992           1993           1994           1995           1996
## 6.051336e-07 -1.369550e-05 -2.092534e-05 7.668013e-07 -5.571115e-06
##           1997           1998           1999           2000           2001
## -9.922250e-06 -2.262969e-06 -2.872153e-06 -1.595609e-06 2.935469e-05
##           2002           2003           2004           2005           2006
## 2.025961e-06 -1.264229e-06 -3.763984e-06 -3.742069e-06 -3.346146e-06
##           2007           2008           2009           2010           2011
## 4.200462e-06 -1.503429e-05 2.331515e-05 -5.152339e-06 -5.596069e-06
##           2012           2013           2014           2015
## -8.750080e-06 -1.380735e-05 -1.056487e-05 -1.350436e-05
```

Practice (if any time remains)

- Compute the standard deviation of every mtcars column.
- Use the following list of formulas to fit multiple models to mtcars

```
formulas <- list(mpg ~ disp,  
                 mpg ~ I(1 / disp),  
                 mpg ~ disp + wt,  
                 mpg ~ I(1 / disp) + wt)
```

- Copy and run the following code to obtain 50 bootstrap samples

```
bootstrap <- function(df) {  
  df[sample(nrow(df), replace = TRUE), , drop = FALSE]  
}  
  
samples <- map(1:50, ~bootstrap(mtcars))
```

- Fit the following model to each bootstrap sample: `mpg ~ disp`
- Extract R^2 and plot the distribution